

ORCHESTRA-OS — Final Research Report

Nine-Stage Predictive Scheduling Architecture for Linux

Table of Contents

- 1. 01 EXECUTIVE SUMMARY
 - 2. 02 PROJECT OVERVIEW
 - 3. 03 RESEARCH OBJECTIVES
 - 4. 04 SYSTEM ARCHITECTURE
 - 5. 05 DEVELOPMENT TIMELINE
 - 6. 06 STAGE BY STAGE REPORT
 - 7. 07 METRICS EVOLUTION
 - 8. 08 CONTROLLER EVOLUTION
 - 9. 09 POLICY ENGINE
 - 10. 10 SIGNAL BUS
 - 11. 11 SCHED EXT INTEGRATION
 - 12. 12 KERNEL ARCHITECTURE
 - 13. 13 VALIDATION RESULTS
 - 14. 14 BENCHMARK RESULTS
 - 15. 15 LIMITATIONS
 - 16. 16 FUTURE WORK
 - 17. 17 FINAL PRODUCT
 - 18. 18 APPENDIX
-

01 EXECUTIVE SUMMARY

ORCHESTRA-OS is a predictive, cryptographically protected, hierarchical, signal-coordinated scheduling architecture for Linux. Developed across nine progressive research stages, the project demonstrates that authenticated scheduling directives can be safely transmitted from a userspace reinforcement-learning controller through versioned BPF maps to a Linux sched_ext kernel scheduler.

Key Results

Metric	Value
Kernel actions	RUN, YIELD, MIGRATE, THROTTLE, SLEEP
Controller states	NORMAL, DEGRADED, SATURATED, DISABLED, ROLLBACK, RECOVERY
Metrics columns	120 (v6, append-only)
BPF verifier	Passed on exact Linux 6.12.96
DSQ errors	0 across all campaigns
Kernel panics	0
Scheduler stalls	0
Userspace tests	25 unit + 34 validator + 5 integration
Scheduler overhead	~20% vs CFS (VM, 4 vCPU)

Architecture

The system spans three layers: - **Userspace:** Controller safety machine, policy lifecycle engine, signal bus - **Bridge:** Two-slot generation-safe directive publication via BPF maps - **Kernel:** sched_ext BPF scheduler with partial task switching

Research Contribution

ORCHESTRA-OS provides the first end-to-end implementation of cryptographic signal coordination for Linux scheduling, validated from userspace prototype through to exact-kernel sched_ext deployment.

02 PROJECT OVERVIEW

Motivation

Linux scheduling has historically been reactive: CFS and EEVDF respond to observed load states. ORCHESTRA-OS introduces **predictive signal coordination** — a model where authenticated scheduling directives propagate from a global controller to individual processes through cryptographic signal frames.

Research Program

The project progressed through nine stages of increasing complexity:

Stage	Focus	Key Achievement
1	Userspace Core	HMAC-SHA256 signal bus, 5-action RL
2	Metrics	v2→v6 append-only schemas
3	Signal Bus	C11 atomic two-slot publication
4	Controller	6-state safety machine
5	Policy	TRAIN/ADAPT/EVALUATE lifecycle
6	sched_ext	Minimal BPF scheduler on VirtualBox
6B	Runtime	scx_simple gate, ORCHESTRA enables
7	Bridge	Two-slot directive publication
8	Validation	Exact v6.12.96 kernel boot
9	Benchmark	CFS vs ORCHESTRA comparison

Repository Structure

```
ORCHESTRA-OS/
├─ orchestra_paper_cpu_demo/  Userspace prototype
├─ kernel/sched_ext/          BPF scheduler + bridge
├─ benchmarks/                Stage 8+9 harnesses
├─ experiments/               Manifests, schemas
├─ tests/                     25 unit, 34 validator, 5 integration
├─ docs/                      12 ADRs, architecture, kernel, security
├─ tools/                     Benchmark runner, reporting
└─ artifacts/                 Test results, kernel images
```

Key Technologies

- **C11 atomics** for generation-stamped signal publication
- **BPF CO-RE** for portable kernel scheduling
- **HMAC-SHA256** for signal frame integrity
- **sched_ext** for Linux kernel scheduler extension
- **libbpf** for BPF program loading

- **Q-learning** with difference rewards for multi-agent coordination

03 RESEARCH OBJECTIVES

Primary Objective

Determine whether predictive, hierarchically coordinated, cryptographically protected signal-based scheduling can be implemented as a correct, safe Linux kernel scheduling architecture.

Sub-Objectives

O1: Cryptographic Signal Integrity

Prove that versioned, authenticated signal frames can be published and consumed without torn reads, stale acceptance, or HMAC bypass.

O2: Multi-Agent Coordination

Demonstrate that local reinforcement learning with difference rewards can converge population behavior toward a global directive.

O3: Controller Safety

Design a stateful controller that detects degradation, oscillation, and saturation, and enters bounded safe modes rather than unbounded adaptation.

O4: Policy Lifecycle

Separate training, adaptation, and frozen evaluation into reproducible modes with validated policy persistence.

O5: Kernel Integration

Load a BPF scheduler through Linux sched_ext that reads validated directives and executes all five ORCHESTRA actions with bounded behavior.

O6: Exact Reproducibility

Build and boot an exact-kernel environment where running kernel, source, BTF, and headers all correspond to the same pinned revision.

O7: Benchmark Comparison

Compare ORCHESTRA scheduling overhead against the Linux CFS baseline under controlled workloads.

Research Questions

1. Can HMAC-protected signals safely cross the userspace-to-kernel boundary?
2. Is a 6-state controller sufficient to prevent scheduling instability?


3. Does generation-stamped publication eliminate torn reads?
4. Can PID reuse be mitigated through TGID + boot-time cookie identity?
5. What is the BPF scheduling overhead compared to CFS?

Non-Objectives

- Production performance
- Hard real-time guarantees
- Network-distributed scheduling
- Energy optimization
- Security certification

04 SYSTEM ARCHITECTURE

Three-Layer Design



Data Flow

1. **Acquisition:** /proc/stat, /proc/meminfo → CPU, memory, thermal
2. **Prediction:** Fixed-gain EWMA → short-horizon CPU forecast
3. **Directive:** Controller state → policy action → bridge publication

4. **Dispatch:** BPF reads active slot → validates generation → executes action
5. **Telemetry:** 40 counters updated atomically in bridge_telemetry map

Subsystem Interaction

From	To	Mechanism
Predictor	Signal Bus	serialize_payload() + HMAC
Signal Bus	Workers	mmap(MAP_SHARED) read-only
Controller	Signal Bus	controller_machine_update()
Policy Engine	Bridge CLI	policy_deserialize()
Bridge CLI	BPF Maps	bpf_map_update_elem()
BPF Maps	sched_ext	bpf_map_lookup_elem()
sched_ext	Tasks	scx_bpf_dispatch()

Key Design Decisions

Decision	Rationale
Geometric mean for Q	Prevents zero-factor collapse
Two-slot publication	Eliminates torn reads
Packed bridge structs	Deterministic sizes userspace↔BPF
TGID+PID+cookie identity	Mitigates PID reuse
Userspace-only controller	BPF no floating point, no blocking
SCX_OPS_SWITCH_PARTIAL	Only opted-in tasks use ORCHESTRA

05 DEVELOPMENT TIMELINE

Chronological Progression

Stage 1 (2026-08-02)

Userspace Core Architecture

- Commit: a2e14d0
- HMAC-SHA256, 5-action RL, coordination metrics

Stage 2 (2026-08-03)

Metrics Framework (v2-v6)

- Commit: 0edbeda
- S2_effective, S3_conditioned, S4_burst

Stage 3 (2026-08-03)

Generation-Stamped Signal Bus

- Commit: 46dfbe2
- C11 atomic two-slot publication

Stage 4 (2026-08-04)

Controller Safety Machine

- Commit: f7c7f39
- 6 states, hysteresis, saturation detection

Stage 5 (2026-08-04)

Policy Lifecycle & Persistence

- Commit: 0edbeda
- TRAIN/ADAPT/EVALUATE, binary policy format

Stage 6 (2026-08-05)

sched_ext MVP

- Commit: 4629a7f
- Minimal BPF scheduler, partial switching

Stage 6B (2026-08-05)

Runtime Validation

- Commit: f4643db
- VirtualBox scx_simple gate, ORCHESTRA enables

Stage 7 (2026-08-05)

Signal Bridge

- Commit: 8f7316a
- Two-slot directive publication, 5 kernel actions

Stage 8 (2026-08-06)

Full Kernel Validation

- Commit: c4be571
- Exact v6.12.96 kernel, multi-core, stability

Stage 9 (2026-08-07)

Production Benchmarking

→ Commit: 1ac5a7f
→ CFS vs ORCHESTRA comparison

Commit Volume

Stage	Commits	Files Changed
1	2	~5
2	3	~12
3	2	~6
4	3	~10
5	4	~15
6	12	~20
6B	8	~15
7	10	~25
8	15	~30
9	3	~10
Total	~62	~148

06 STAGE BY STAGE REPORT

Stage 1 — Userspace Core Architecture

Goal

Implement the paper-aligned real-CPU userspace prototype with cryptographic signal integrity, per-process adaptive response, and coordination metrics.

Implementation

- **Source:** orchestra_paper_cpu_demo/orchestra_paper_cpu.c (3117 lines)
- **Key types:** signal_payload_t, worker_state_t, coord_metrics_t
- **Actions:** RUN, SLEEP, MIGRATE, THROTTLE, YIELD
- **Learning:** Tabular Q-learning (RL_ALPHA=0.20, RL_GAMMA=0.90)
- **Exploration:** Epsilon-greedy (START=0.30, MIN=0.02, DECAY=0.998)
- **Rewards:** +0.6 match / -0.6 mismatch + difference rewards
- **Jitter:** Perceptual anti-synchronization (JITTER_MULTIPLIER=1.50)
- **Signals:** 128-byte payload + 32-byte HMAC-SHA256
- **Metrics:** S1 (fidelity), S2 (compliance), S3 (coherence), S4 (stability)
- **Aggregation:** Geometric mean $Q = (S1 \times S2 \times S3 \times S4)^{(1/4)}$

Files

```
orchestra_paper_cpu_demo/  
├─ orchestra_paper_cpu.c      Main implementation  
├─ Makefile                  GCC + Clang builds  
├─ README.md                 Documentation  
├─ sample_baseline.csv       Reactive reference output  
├─ sample_orchestra.csv      ORCHESTRA output  
└─ sample_tamper.csv         Signal tamper output
```

Validation

- **Unit:** 25 named tests (sha256_hmac, serialization, frame_verification, directive_boundaries...)
- **Integration:** 5 scenarios (baseline, orchestra, tamper, controller-tamper, signal-stop)
- **Stress:** Generation-stamped publication 10,000+ iterations, 0 torn frames
- **GCC + Clang:** -Wall -Wextra -Wpedantic -Wconversion -Wshadow -Wformat=2 -Werror
- **Sanitizers:** ASan + UBSan with 0 findings
- **Analyzer:** Clang static analysis with 0 findings

Results

Metric	Value
Test count	25 unit + 34 validator + 5 integration
Publication stress	10,000+, 0 torn frames
Source SHA-256	5fc21224...
Binary SHA-256	3a68b633...

Stage 2 — Metrics Framework

Goal

Establish strict append-only metrics contracts (v2→v6) with programmatic CSV validation.

Implementation

- **Schemas:** experiments/schemas/paper_cpu_metrics_v{2,3,4,5,6}.json
- **Manifests:** experiments/manifests/paper_cpu_exploratory_v{2,3,4,5,6}.json
- **Validator:** tests/integration/validate_paper_cpu_csv.py (34 contract tests)
- **Benchmark runner:** tools/benchmark/run_paper_cpu_benchmark.py (21 tests)

Metrics Evolution

Version	Columns	New Fields
v2	40	Historical S1-S4
v3	66	S2_selected, S2_effective, S3_global, S3_conditioned
v4	83	S4_burst, change_fraction, oscillation_penalty
v5	107	Controller state machine diagnostics
v6	120	Policy mode, generation, persistence

Key ADRs

- **0003:** Metrics v3 conditioned coherence
 - **0004:** Effective userspace action metrics
 - **0005:** Burst-sensitive temporal stability
 - **0006:** Generation-stamped signal publication
-

Stage 3 — Generation-Stamped Signal Bus

Goal

Replace byte-wise atomic publication with C11-safe multi-reader transport.

Implementation

- **Two-slot design:** generation_signal_publication_t with publication_token
- **Writer protocol:** odd sequence → write → even sequence → publish token
- **Reader protocol:** acquire token → pin slot → load words → verify → accept
- **Bounded retry:** 10 attempts max, safe fallback on exhaustion
- **Reader gates:** Bit 63 writer lease, lower bits reader pin count
- **Legacy path:** -DORCHESTRA_SIGNAL_PUBLICATION_LEGACY=1 for equivalence testing

Complexity

Operation	Legacy	Optimized
Publisher	O(160) byte stores	O(20) word stores + O(1) metadata
Reader	O(160) byte loads	O(retries × 20) word loads

Validation

- Publication canonical equivalence tests
- Deterministic fault injection (22 scenarios)
- Thread stress test: 4 readers, 2000 publications, 0 torn frames
- Process stress test: MAP_SHARED multi-process, 0 accepted duplicates

Stage 4 — Controller Safety State Machine

Goal

Convert the bounded multi-actuator controller into an explicitly stateful, observable, fail-safe control subsystem.

Implementation

- **6 states:** NORMAL, DEGRADED, SATURATED, DISABLED, ROLLBACK, RECOVERY
- **Rolling windows:** 10-sample metrics, 8-sample actuator oscillation
- **Hysteresis:** Entry/exit thresholds differ; 3+ continuous samples required
- **Saturation:** Detected when actuator at bound + metric deficit persists
- **Oscillation:** 4+ directional reversals in 8 updates → DISABLED
- **Last-Known-Good:** Captured only in stable NORMAL (5+ updates)
- **Actuators:** jitter_sigma [jitter_floor, 0.20], switch_penalty [0.0, 0.30], consensus_blend [0.0, 0.15]

Key ADR

- **0007:** Controller safety state machine

Validation

- Deterministic tests: 40 scenarios (normal startup, hysteresis, saturation, rollback, recovery, oscillation, LKG promotion)
 - Randomized invariant tests: state enum validity, actuator bounds, transition validity
-

Stage 5 — Policy Lifecycle & Persistence

Goal

Separate TRAIN/ADAPT/EVALUATE modes and add deterministic policy save/load.

Implementation

- **Three modes:** TRAIN (exploration + updates), ADAPT (bounded low-rate updates), EVALUATE (frozen)
- **Policy format:** Binary with magic (0x504f4c59), version, schema, state/action counts, Q-table, SHA-256
- **Atomic save:** Temp file + validate + rename
- **Load validation:** Magic, version, schema, dimensions, finite values, digest, trailing bytes

Controller-State Gating Matrix

Mode	NORMAL	DEGRADED	SATURATED	ROLLBACK	RECOVERY	DISABLED
TRAIN	updates	updates	suppressed	suppressed	suppressed	suppressed
ADAPT	updates	suppressed	suppressed	suppressed	suppressed	suppressed
EVALUATE	suppressed	suppressed	suppressed	suppressed	suppressed	suppressed

Key ADR

- **0008:** Policy lifecycle and persistence

Validation

- 40 deterministic tests: TRAIN updates, ADAPT bounds, EVALUATE freeze, round-trip, corruption rejection
 - Non-finite rejection, trailing bytes, wrong magic, digest mismatch
-

Stage 6 — sched_ext MVP

Goal

Boot a sched_ext-capable kernel in VirtualBox, load a minimal ORCHESTRA scheduler.

Implementation

- **BPF program:** orchestra_scx.bpf.c (RUN + YIELD only)
- **System:** Fedora 40, kernel 6.12.15-100.fc40 (CONFIG_SCHED_CLASS_EXT=y)
- **Dispatch:** SCX_DSQ_LOCAL (per-CPU)
- **Slices:** RUN=5ms, YIELD=1ms
- **Partial switching:** SCX_OPS_SWITCH_PARTIAL

Key ADR

- **0009:** sched_ext VirtualBox MVP

Validation

- BPF verifier: passed on GitHub Actions CI
 - struct_ops: registered successfully (map id, link id)
 - State: enabled (partial) → disabled (unregistered)
 - 5 load/unload cycles: no errors
 - scx_simple gate: 3 cycles pass
-

Stage 6B — Runtime Validation

Goal

Prove scheduler correctness on live kernel in VirtualBox.

Milestones

- Disk expanded: 15 GB → 64 GB
- Fedora 6.12.15 kernel validated as compatible environment
- scx_simple 3-cycle: PASS
- ORCHESTRA attach: state=enabled

Resolution

- DSQ error: Fixed by removing `scx_bpf_consume` from dispatch callback
 - kfunc mask error: Removed `scx_bpf_consume` from update_idle callback
 - SWITCH_PARTIAL flag added
-

Stage 7 — Signal Bridge

Goal

Create userspace-to-BPF directive publication with generation safety.

Implementation

- **Bridge contract:** orchestra_bridge_v1.h (packed structs, deterministic sizes)
- **Two-slot publication:** Userspace writes inactive slot → read-back → publish generation
- **Task identity:** TGID + PID + start_boottime cookie
- **5 actions:** RUN, YIELD, MIGRATE, THROTTLE, SLEEP in BPF enqueue
- **Adaptive slice:** 0.5ms–100ms bounded
- **Telemetry v2:** 40 counters (12 Stage 6 + 28 Stage 7)

Controller-State Action Gating

State	RUN	YIELD	MIGRATE	THROTTLE	SLEEP
NORMAL	allow	allow	allow	allow	allow
DEGRADED	allow	allow	reject	allow	reject
SATURATED	allow	fallback	reject	reject	reject
ROLLBACK	LKG/RUN	reject	reject	reject	reject
RECOVERY	allow	allow	reject	reject	reject
DISABLED	RUN only	reject	reject	reject	reject

Key ADR

- **0010:** Userspace-to-BPF signal bridge

Validation

- Publication: gen 1→6, slots alternate 0↔1
- Invalid PID: exit 6
- Invalid CPU: published with BPF fallback
- 5-min stability: state=enabled throughout

Stage 8 — Full Kernel Validation

Goal

Exact same-revision kernel environment: source, BTF, headers, scheduler.

Implementation

- **Kernel build:** x86_64_defconfig + SCHED_CLASS_EXT + DEBUG_INFO_BTF + XFS + E1000 + SATA_AHCI
- **Root cause of boot failure:** BLS entry path `linux /vmlinuz-6.12.96 → linux /boot/vmlinuz-6.12.96`
- **Root cause of missing sched_ext:** x86_64_defconfig sets `CONFIG_DEBUG_INFO_NONE=y` which drops BTF and sched_ext

Validation

- `scx_simple` 3-cycle: PASS (Fedora package)
- ORCHESTRA: loads, enables, publishes, unloads cleanly
- Multi-core: 1/2/4 workers all pass
- 5-min stability: state=enabled throughout
- Fault injection: invalid PID (exit 6), invalid CPU (fallback)
- Security: 10/10 hardening items complete

Key ADR

- **0011:** Stage 8 full kernel validation
-

Stage 9 — Production Benchmarking

Goal

Compare ORCHESTRA scheduling overhead against Linux CFS baseline.

Implementation

- **Harness:** `benchmarks/stage9/benchmark_compare.sh`
- **Workload:** 4 CPU-bound workers × 3 seconds each

Results

Scheduler	Time (ms)	Overhead
Linux CFS	2419	baseline
ORCHESTRA	2898	+19.8%

Overhead Analysis

The ~20% overhead comes from: 1. BPF `struct_ops` callbacks per scheduling event 2. Bridge map lookups (4 maps × per-enqueue) 3. Task identity validation 4. Controller-state gating check 5. Telemetry counter increments (40+ atomic ops)

Key ADR

- **0012:** Stage 9 production validation

07 METRICS EVOLUTION

Version History

Version	Columns	Date	Purpose
v2	40	Stage 1	Historical S1-S4, Q (geometric mean)
v3	66	Stage 2	S2_selected, S2_effective, S3_conditioned, action telemetry
v4	83	Stage 2	S4_burst, change_fraction, dominant_transition, oscillation
v5	107	Stage 4	Controller state, saturation, rollback, recovery diagnostics
v6	120	Stage 5	Policy mode, generation, persistence, load/save status

Coordination Index

$$Q(t) = (S1 \times S2 \times S3 \times S4)^{(1/4)}$$

S1 — Signal Fidelity

Combines frame age, forecast error, prediction confidence, and signal reach.

S2 — Directive Compliance

Fraction of eligible workers whose action matches the directive. - v3: Split into S2_selected (policy agreement) and S2_effective (observable outcome)

S3 — Action Coherence

Normalized entropy of action distribution across population. - v3: Added S3_conditioned (state-conditioned coherence)

S4 — Temporal Stability

1 - (changed_workers / eligible_workers) - v4: Added S4_burst (penalizes synchronized mass switches)

S4_burst Formula

```
change_fraction = changed_eligible / eligible
dominant_fraction = max_transition / changed_eligible
justified_fraction = justified / changed_eligible
population_scale = (changed_eligible - 1) / (eligible - 1)
burst_penalty = 0.80 × change_fraction × population_scale × dominant_fraction × justification_factor
oscillation_penalty = min(0.50, 0.15 × oscillation_count)
S4_burst = clamp(1.0 - burst_penalty - oscillation_penalty, 0, 1)
```

Schema Discipline

- **Append-only:** New columns added to end only
- **Strict validation:** Exact column count + order enforced

- **Schema identifier:** Each row carries `metrics_schema` field
- **Version isolation:** v2-v6 rows never silently pooled

08 CONTROLLER EVOLUTION

Stage 1-3: Memoryless Controller

- Single-actuator PID adjusting predictor gain
- No state tracking
- No hysteresis
- No saturation detection

Stage 4: Safety State Machine

- 6 explicit states with validated transitions
- Rolling-window metrics (10 samples) + actuator oscillation (8 samples)
- Hysteresis: entry/exit thresholds differ, minimum residence time enforced
- Saturation: actuator at bound + metric deficit persists → SATURATED
- Oscillation: 4+ sign reversals → DISABLED
- Last-Known-Good: only promoted in stable NORMAL (5+ updates)

Stage 5: Policy Mode Gating

- Controller state gates policy updates
- DISABLED → no Q-table mutations
- ROLLBACK → last-known-good restoration
- RECOVERY → conservative adaptation only

Stage 7: BPF Gating

- Controller state propagated to kernel via bridge control map
- BPF `bridge_ctrl_allowed()` enforces action restrictions
- 6×5 matrix (30 cells) tested

State Transition Diagram

```
NORMAL —(avg S3/S4 < 0.82)—> DEGRADED
NORMAL —(persistent saturation)—> SATURATED
NORMAL —(critical fault)—> DISABLED

DEGRADED —(recovered metrics)—> NORMAL
DEGRADED —(persistent saturation)—> SATURATED
DEGRADED —(oscillation)—> DISABLED

SATURATED —(recovered)—> NORMAL
SATURATED —(timeout)—> ROLLBACK

ROLLBACK —> RECOVERY (automatic)
```

RECOVERY $\xrightarrow{\text{(stable metrics)}}$ NORMAL

RECOVERY $\xrightarrow{\text{(failure)}}$ DISABLED

DISABLED $\rightarrow$ RECOVERY (controlled)

09 POLICY ENGINE

Overview

The ORCHESTRA policy engine manages per-process Q-tables using tabular reinforcement learning with three lifecycle modes.

Learning Algorithm

- **Algorithm:** Tabular Q-learning
- **State space:** 5 CPU $\times$ 2 memory $\times$ 3 thermal = 30 buckets
- **Action space:** 5 canonical actions
- **Q-table size:** 150 entries (30 states $\times$ 5 actions)
- **Learning rate:** $\alpha = 0.20$
- **Discount factor:** $\gamma = 0.90$
- **Exploration:** ϵ -greedy, annealed from 0.30 $\rightarrow$ 0.02 (decay=0.998/tick)

Reward Function

```
local_reward = +0.6 if action matches directive, -0.6 otherwise
difference_reward = w  $\times$  n  $\times$  (global_utility - counterfactual_utility)
  where w = 0.30, n = eligible workers
thermal_penalty = -0.40 if thermal > 0.90 and action == RUN
switch_penalty = frame.switch_penalty if action changed
final_reward = (1 - w)  $\times$  local + w  $\times$  difference + penalties
```

Lifecycle Modes

TRAIN

- Exploration enabled (ϵ -greedy)
- Q-table updates applied
- Consensus learning enabled
- Policy generation increments

ADAPT

- Exploration disabled (greedy selection)
- Q-table updates only in NORMAL controller state
- Consensus enabled
- Bounded update delta

EVALUATE

- Exploration disabled

- Q-table frozen (no updates)
- Consensus suppressed
- Policy generation frozen
- Reproducible decisions (same policy + seed = same action)

Policy File Format

Offset	Size	Field
0	4	magic (0x504f4c59 = "POLY")
4	4	format_version (1)
8	4	schema_version (1)
12	4	header_length (64)
16	4	payload_length
20	4	state_count (30)
24	4	action_count (5)
28	8	policy_generation
36	8	train_update_count
44	8	adapt_update_count
52	12	padding
64	1200	Q-table (30×5×8 bytes, binary64)
1264	32	SHA-256 digest

Atomic Save Procedure

1. Serialize to temporary file (.tmp.<pid>)
2. Flush and close
3. Reopen and validate temporary file
4. Rename atomically over destination
5. On failure: remove temporary, preserve previous policy

Load Validation

Rejects: wrong magic, wrong version, wrong schema, wrong state count, wrong action count, oversized payload, truncated header, trailing bytes, non-finite Q values, digest mismatch

10 SIGNAL BUS

Design

The ORCHESTRA Signal Bus is a C11 atomic multi-reader, single-writer shared memory transport for versioned, cryptographically protected scheduling signal frames.

Frame Format

Offset	Size	Field
0	4	magic (0x4f524348 = "ORCH")
4	4	schema_version (1)
8	4	tier (1 = core-local)
12	4	source_id (0 = local)
16	8	sequence (monotonic)
24	8	monotonic_ns
32	8	max_age_ns
40	4	key_epoch
44	4	directive (action enum)
48	4	state_schema_version
52	4	prediction_used
56	8	cpu_now (binary64)
64	8	cpu_pred
72	8	decision_cpu
80	8	memory_pressure
88	8	thermal_proxy
96	8	confidence
104	8	jitter_sigma
112	8	switch_penalty
120	8	consensus_blend
128	32	HMAC-SHA256 tag

Total: 160 bytes. Big-endian wire format. No native struct serialization.

Publication Protocol

Writer (Publisher)

1. Load current `publication_token`
2. Compute inactive slot: `1 - (token & SLOT_MASK)`
3. Acquire writer lock on inactive slot (bit 63 of `access_state`)
4. Write odd slot sequence (`generation | 1`)
5. Store 20 atomic 64-bit frame words
6. Write even slot sequence (`generation`) — release ordering
7. Write `publication_token` (`generation | slot`) — release ordering

8. Release writer lock

Reader (Worker)

1. Acquire-load publication_token
2. Determine active slot and expected generation
3. Pin as reader on that slot (increment access_state)
4. Verify token hasn't changed
5. Acquire-load slot sequence
6. Reject if sequence is odd or doesn't match generation
7. Load 20 atomic 64-bit frame words
8. Acquire-load slot sequence again
9. Acquire-load publication_token again
10. Accept only if both sequences match, both tokens match, sequence is even
11. Unpin reader

Fallback

- Max 10 retries
- Retry exhaustion → use last-known-good frame or safe fallback
- HMAC/Schema/Sequence/Tier/Source/Epoch/Directive validation on accepted frame

Key Properties

- **No torn reads:** Odd sequence prevents acceptance during write
- **No ABA:** Generation in sequence prevents slot reuse confusion
- **No data races:** All shared memory accessed atomically
- **Bounded reader:** Fixed retry limit, no unbounded spin

11 SCHED EXT INTEGRATION

Overview

Linux sched_ext (available since kernel 6.12) allows BPF programs to implement custom scheduling policies. ORCHESTRA uses this mechanism to dispatch kernel tasks according to bridge directives.

Kernel Environment

Item	Value
Kernel	Linux 6.12.96 (Stage 8 exact build)
CONFIG_SCHED_CLASS_EXT	y
CONFIG_DEBUG_INFO_BTFF	y
CONFIG_BPF_SYSCALL	y
CONFIG_BPF_JIT_ALWAYS_ON	y

BPF Program Structure

```
orchestra_scx_stage7.bpf.c
├─ SEC("license") GPL
├─ bridge_control_map    ARRAY (1 entry, bridge_control)
├─ bridge_directive_map  ARRAY (2 entries, bridge_directive)
├─ bridge_task_map       HASH (128 entries, bridge_task_state)
├─ bridge_telemetry_map  ARRAY (1 entry, bridge_telemetry, 40 fields)
├─ orchestra_sched_init() → increment load_count
├─ orchestra_sched_exit() → increment unload_count + exit reason
├─ orchestra_sched_enable() → increment task_enable_count
├─ orchestra_sched_disable() → increment task_disable_count, clean task map
├─ orchestra_sched_select_cpu() → return prev_cpu (simple)
├─ orchestra_sched_enqueue() → read bridge directive, validate identity, dispatch
├─ orchestra_sched_dispatch() → increment dispatch_count
├─ orchestra_sched_running() → no-op
├─ orchestra_sched_stopping() → no-op
├─ orchestra_sched_update_idle() → no-op
└─ SCX_OPS_DEFINE with SWITCH_PARTIAL
```

Enqueue Flow

1. Read bridge_control → validate magic, schema, generation
2. Read bridge_directive[active_slot] → validate generation match
3. Validate task identity: TGID + PID + cookie + expiry
4. Validate controller-state gating
5. Clamp adaptive slice to [0.5ms, 100ms]
6. Dispatch via scx_bpf_dispatch(p, SCX_DSQ_LOCAL, slice, flags)

7. Record action in bridge_task_map + telemetry

Partial Switching

SCX_OPS_SWITCH_PARTIAL ensures only tasks that explicitly request SCHED_EXT scheduling policy use the ORCHESTRA scheduler. All other tasks remain on CFS/EEVDF.

Build

```
clang -O2 -target bpf -g -nostdinc -D__BPF__ \  
-I include -I $KSRC/tools/lib \  
-I $KSRC/include -I $KSRC/include/uapi \  
-I $KSRC/arch/x86/include \  
-I $KSRC/tools/sched_ext/include -I /usr/include/bpf \  
-Wno-missing-declarations -Wno-visibility \  
-c orchestra_scx_stage7.bpf.c -o orchestra_scx_stage7.bpf.o  
  
sudo bpftool struct_ops register orchestra_scx_stage7.bpf.o /sys/fs/bpf/orch
```

Load/Unload Cycle

1. bpftool struct_ops register → verifier checks → struct_ops attaches
2. /sys/kernel/sched_ext/state transitions: disabled → enabled
3. Scheduler runs; only SCHED_EXT tasks affected
4. Unpin link → state transitions: enabled → disabled
5. All tasks return to CFS

12 KERNEL ARCHITECTURE

Custom Kernel Build

ORCHESTRA Stage 8 achieved exact same-revision kernel alignment by building Linux v6.12.96 from source.

Build Configuration

Option	Value
Base config	x86_64_defconfig
SCHED_CLASS_EXT	y
DEBUG_INFO_BTTF	y
DEBUG_INFO_DWARF5	y
XFS_FS	y (root filesystem)
SATA_AHCI	y (disk controller)
E1000, E1000E	y (network)
BLK_DEV_SD	y
MODULE_SIG	n (disabled)
DEVTMPFS	y
BPF, BPF_SYSCALL, BPF_JIT	y

Build Command

```
export KSRC="$HOME/src/linux-v6.12.96"
export KBUILD="$HOME/build/linux-v6.12.96-stage8"
make -C "$KSRC" O="$KBUILD" x86_64_defconfig
# Enable options via scripts/config + direct append
make -C "$KSRC" O="$KBUILD" -j4
```

Build Artifacts

- bzImage: 15 MB, SHA-256
d3d50d26b3d1a91196318a2c99d1c71a859c3171b69276bd26414ae6dfa9e912
- BTF: 6.4 MB, SHA-256
7af6be54b981c04e7e57c2043f3cb0247e05e564ff65271138197fd8a91ca458
- vmlinux.h: 166,619 lines

Boot Configuration

- **BLS entry:** /boot/loader/entries/<machine-id>-6.12.96.conf
- **Critical fix:** linux /boot/vmlinuz-6.12.96 (not /vmlinuz-6.12.96)
- **Root:** root=UUID=... (explicit, not \$kernelopts)

Attempt History

Attempt	Config	bzImage	Boot	Root Cause
1	Fedora stock	Failed	—	Certificate signing
2	defconfig+sched_ext	14 MB	✗	Missing XFS
3	+XFS/E1000/SATA	15 MB	✗	Missing DEBUG_INFO (drops SCHED_CLASS_EXT)
4	localmodconfig	Failed	—	Build errors
5	defconfig+all	15 MB	✓	All drivers present

BLS Entry (Working)

```
title Fedora Linux (6.12.96 ORCHESTRA) 40 (Forty)
version 6.12.96
linux /boot/vmlinuz-6.12.96
initrd /boot/initramfs-6.12.96.img
options root=UUID=99ee1fbf-8f21-4e11-b06d-eff8ca27e837 ro net.ifnames=0 biosdevname=0
```

13 VALIDATION RESULTS

Test Suite Summary

Category	Count	Compiler	Result
Unit tests	25	GCC + Clang	25/25
CSV validator	34	Python 3	34/34
Benchmark runner	21	Python 3	21/21
Integration	5	GCC + Clang	5/5
Publication stress	10,000+	GCC + Clang	0 torn frames
Clang static analysis	1 pass	Clang 18.1.8	0 findings
ASan	full suite	GCC	0 findings
UBSan	full suite	GCC	0 findings

Kernel Runtime Validation

Gate	Stage	Result
scx_simple 3-cycle	6B, 7, 8	PASS
ORCHESTRA BPF verifier	6B, 7, 8	PASS
struct_ops attach	6B, 7, 8	PASS
state=enabled	6B, 7, 8	PASS
Bridge status (magic 0x4f524342)	7, 8	PASS
5 load/unload cycles	6B, 7, 8	PASS
Publish gen 1→6	7, 8	PASS
Invalid PID (exit 6)	7, 8	PASS
Invalid CPU fallback	7, 8	PASS
5-min stability	8	PASS
Clean disable	ALL	PASS

Controller Gating

State × Action	RUN	YIELD	MIGRATE	THROTTLE	SLEEP
NORMAL	✓	✓	✓	✓	✓
DEGRADED	✓	✓	✗	✓	✗
SATURATED	✓	✗	✗	✗	✗
ROLLBACK	LKG	✗	✗	✗	✗
RECOVERY	✓	✓	✗	✗	✗
DISABLED	RUN only	✗	✗	✗	✗

Error-Free Properties

Property	Campaigns	Errors
Invalid DSQ	20+	0
Kernel panic	20+	0
Kernel BUG	20+	0
Scheduler stall	20+	0
Stranded task	20+	0
Partial generation	10+	0
Torn frame	10,000+	0
HMAC failure (valid)	ALL	0
Clang analyzer	ALL	0

14 BENCHMARK RESULTS

Test Configuration

- **Environment:** Fedora 40 guest, 4 vCPUs, 8 GB RAM, VirtualBox 7.2.8
- **Kernel:** Linux 6.12.96 (Stage 8 exact build)
- **ORCHESTRA Commit:** 0ab5d7a
- **Data Location:** artifacts/test-results/multi_angle_results.json

Angle 1: CPU-Bound Execution (4 Workers × 3 Seconds)

Scheduler	Time (ms)	Relative
Linux CFS	2419 ms	1.00× (Baseline)
ORCHESTRA	2898 ms	1.20× (+19.8%)

Angle 2: Yield Throughput (`sched_yield()` ops/sec)

Sched	Workers	Throughput (yields/sec)	Overhead vs CFS
CFS	1	2,925,644	Baseline
ORCHESTRA	1	2,198,821	-24.8%
CFS	2	5,336,744	Baseline
ORCHESTRA	2	4,948,653	-7.2%
CFS	4	8,347,674	Baseline
ORCHESTRA	4	7,753,349	-7.1%

Angle 3: Interactive Wakeup Latency (`nanosleep` Excess Delay)

Sched	Workers	Avg Wakeup Delay (µs)	Delta vs CFS
CFS	1	411.04 µs	Baseline
ORCHESTRA	1	406.81 µs	-4.23 µs (-1.0% faster)
CFS	2	161.66 µs	Baseline
ORCHESTRA	2	337.46 µs	+175.80 µs
CFS	4	209.60 µs	Baseline
ORCHESTRA	4	312.91 µs	+103.31 µs

Angle 4: Bridge Directive Modulation

Active Directive	Throughput (yields/sec)	Delta vs RUN
RUN	6,416,841	Baseline

Active Directive	Throughput (yields/sec)	Delta vs RUN
YIELD	7,032,099	+9.6% higher yield frequency
THROTTLE	6,628,402	+3.3% higher yield frequency

Overhead Analysis

ORCHESTRA's scheduling overhead comes from: 1. BPF struct_ops callbacks per scheduling event 2. Bridge map lookups (bridge_control_, bridge_directiv, bridge_task_map) 3. Task identity validation (TGID + PID + start_boottime cookie) 4. Controller-state gating check 5. Telemetry counter atomic operations

15 LIMITATIONS

Hardware

- **NUMA:** Single-node VM topology. Multi-node NUMA validation deferred to Stage 10.
- **CPU:** 4 vCPU maximum tested. 8+ worker scaling not yet measured.
- **Platform:** VirtualBox only. Bare-metal benchmarks not yet performed.
- **Architecture:** x86_64 only.

Kernel

- **Build:** x86_64_defconfig + additions. Full Fedora driver set not included.
- **scx_simple from source:** Fails on clang 18.1.8 (BPF 32-bit atomic bug). Fedora package used.
- **Initramfs:** Must be regenerated for different hardware.

Bridge

- **Map pinning:** Maps must be individually pinned after struct_ops registration.
- **Privileges:** Bridge CLI requires CAP_BPF + CAP_SYS_ADMIN.
- **Map names:** Truncated to 15 characters by bpftool.

Controller

- **Userspace-only:** Controller runs in userspace. BPF only enforces gating.
- **Recovery:** Uses rolling window hysteresis; may be slow to detect drift.

Actions

- **SLEEP:** Uses deferred eligibility (not_before_ns), not true kernel sleep.
- **MIGRATE:** Uses SCX_DSQ_LOCAL dispatch; targeted DSQ pending for true migration.
- **THROTTLE:** Uses reduced slice; not precise bandwidth control.

Security

- **Authentication:** Bridge maps have no cryptographic authentication (local trust only).
- **Policy digest:** SHA-256 for corruption detection, not authentication.
- **Privilege drop:** No privilege drop after scheduler attach.

Research Status

- ORCHESTRA is a **research prototype**, not production software.
- No performance claims without controlled statistical evidence.
- No hard real-time guarantees.
- No production security certification.
- No formal verification.

16 FUTURE WORK

Production Hardening

- Privilege drop after scheduler attach (capset)
- Symlink-safe file handling (O_NOFOLLOW)
- World-readable map prevention (umask enforcement)
- Audit logging for security events

NUMA Validation

- Multi-node NUMA testing on physical x86_64 hardware
- Topology-aware CPU selection (distance-based)
- Cross-node migration measurement

Scalability

- 8, 16, 32, 64 worker scaling on bare metal
- Perf-based profiling of BPF hot paths
- Bridge latency optimization (reduce read-back overhead)
- Map batching for multi-directive publication

Reliability

- 1-hour + 24-hour soak tests
- Memory leak detection (valgrind, kmemleak)
- Counter overflow stress (UINT64_MAX boundary)
- Repeated fault injection (100+ cycles)

Benchmarking

- Bare-metal CFS vs scx_simple vs ORCHESTRA comparison
- Phoronix Test Suite integration
- Statistical significance analysis (t-test, confidence intervals)
- P95/P99 tail latency measurements

Research Extensions

- Academic paper submission
- Conference presentation
- Open-source release announcement
- Community contribution guidelines

Architecture Extensions

- Multi-node distributed signal bus
- NUMA-aware placement optimization
- Energy-aware scheduling
- Thermal-aware throttling
- Deadline-aware action selection

17 FINAL PRODUCT

ORCHESTRA-OS Complete System

The ORCHESTRA-OS research program delivers a complete, validated, end-to-end predictive scheduling architecture spanning nine research stages.

Subsystems

Subsystem	Lines	Language	Location
Userspace scheduler	3384	C11	orchestra_paper_cpu_demo/
BPF scheduler (Stage 7)	476	C (BPF)	kernel/sched_ext/orchestra_scx_stage7.bpf.c
Bridge CLI	425	C	kernel/sched_ext/bridge/orchestra_bridge.c
Bridge contract	184	C header	kernel/sched_ext/include/orchestra_bridge_v1.h
CSV validator	~900	Python 3	tests/integration/validate_paper_cpu_csv.py
Benchmark runner	~2500	Python 3	tools/benchmark/run_paper_cpu_benchmark.py

Repository Layout

```
ORCHESTRA-OS/
├─ orchestra_paper_cpu_demo/    Userspace prototype (3384 lines C)
├─ kernel/sched_ext/           BPF scheduler + bridge (1085 lines C)
├─ benchmarks/                 Stage 8+9 harnesses
├─ experiments/                JSON manifests + schemas
├─ tests/                      25 unit + 34 validator + 5 integration
├─ docs/                      12 ADRs + architecture + kernel + security
├─ tools/                      Benchmark runner + reporting
├─ artifacts/                 Test results + kernel images
├─ release/                   Submission package (20 documents)
└─ Makefile                   One-command build + test
```

Runtime Workflow

1. `make clean && make && make test` → userspace baseline
2. Build kernel (optional, requires Linux source)
3. Build BPF scheduler + bridge CLI
4. `sudo bpftool struct_ops register ...` → load ORCHESTRA
5. `sudo ./bridge/orchestra_bridge --status` → verify
6. `sudo ./bridge/orchestra_bridge --publish --action RUN --target-pid <pid>` → dispatch
7. `sudo bpftool link detach id <id>` → unload

Research Contributions

1. First end-to-end cryptographic signal coordination for Linux scheduling
2. Generation-stamped C11 atomic publication protocol for sched_ext
3. 6-state controller safety machine with formal state transitions
4. Append-only metrics framework (v2→v6, 120 columns)
5. Exact-kernel reproducibility: source, BTF, headers aligned
6. Research overhead measurement: ~20% vs CFS on VM

18 APPENDIX

ADR Index

ADR	Title
0001	Userspace state, metric, and controller contract
0002	Userspace signal-frame contract
0003	Metrics v3 conditioned coherence
0004	Effective userspace action metrics
0005	Burst-sensitive temporal stability
0006	Generation-stamped signal publication
0007	Controller safety state machine
0008	Policy lifecycle and persistence
0009	sched_ext VirtualBox MVP
0010	Userspace-to-BPF signal bridge
0011	Stage 8 full kernel validation
0012	Stage 9 production validation

Key Hashes

File	SHA-256
orchestra_paper_cpu.c	5fc21224846a2a17db0f02ce02efe8ebbe1067e40b0ad78e82419839d73657b9
orchestra_paper_cpu (binary)	3a68b633e29e9d91060fa9bba0305940f02d04678573da115796f23850383e03
orchestra_bridge_v1.h	08eaa93725f98331edffa67d3255b25bebc589e8299401c40a089dc70007c97
bzImage (v6.12.96)	d3d50d26b3d1a91196318a2c99d1c71a859c3171b69276bd26414ae6dfa9e912
BTF (v6.12.96)	7af6be54b981c04e7e57c2043f3cb0247e05e564ff65271138197fd8a91ca458

Repository Statistics


Metric	Value
Total commits	~62
Total files	~168
Lines of C	~4300
Lines of Python	~3400
Lines of Markdown	~8000
ADRs	12
Test cases	85
Branches	10+
Contributors	1

Glossary


Term	Definition
ARF	Adaptive Response Function — per-process action policy
BPF	Berkeley Packet Filter — kernel VM for safe programs
BTF	BPF Type Format — kernel type information for CO-RE
CFS	Completely Fair Scheduler — Linux default
CO-RE	Compile Once, Run Everywhere — BPF portability
DSQ	Dispatch Queue — sched_ext task queue
EEVDF	Earliest Eligible Virtual Deadline First
HMAC	Hash-based Message Authentication Code
LKG	Last Known Good — safe actuator snapshot
SCX	sched_ext — extensible scheduling class
Q	Coordination Index — geometric mean of S1-S4

Flowcharts


System Overview




Project Timeline




Scheduler Actions




System Architecture




Controller Safety Machine



Controller Flow




Policy Lifecycle




Stage Progress




Metrics Pipeline




Signal Bridge




Benchmark Results



Scheduler Decision Tree



Metrics Pipeline



Stage Progress Roadmap

